

UNIVERSITY OF PENNSYLVANIA

ESE 546: PRINCIPLES OF DEEP LEARNING

HOMEWORK 2

1 Read the following instructions carefully before beginning to work on the homework.

- 2 • You will submit solutions typeset in L^AT_EX on Gradescope (strongly encouraged). You can
3 use hw_template.tex on Canvas in the “Homeworks” folder to do so. If your handwriting is
4 unambiguously legible, you can submit PDF scans/tablet-created PDFs.
5 • Clearly indicate the name and Penn email ID of all your collaborators on your submitted
6 solutions.
7 • Start a new problem on a fresh page and mark all the pages corresponding to each problem.
8 Failure to do so may result in your work not graded completely.
9 • For each problem in the homework, you should mention the total amount of time you spent
10 on it. This helps us keep track of which problems most students are finding difficult.
11 • You can be informal while typesetting the solutions, e.g., if you want to draw a picture feel
12 free to draw it on paper clearly, click a picture and include it in your solution. Do not spend
13 undue time on typesetting solutions.
14 • You will see an entry of the form “HW 1 PDF” where you will upload the PDF of your
15 solutions. You will also see entries like “HW 1 Problem 1 Code” and “HW 1 Problem 3
16 Code” where you will upload your solution for the respective problems.
17 • **For each programming problem/sub-problem, you should create a fresh .py file.** This
18 file should contain **all** the code to reproduce the results of the problem/sub-problem, e.g., it
19 should save the plot that is required (correctly with all the axes, title and legend) as a PDF
20 in the same directory. You will upload the .py file as your solution for “HW 1 Problem 3
21 Code” or “HW 1 Problem 3 Code”. Name your file as pennkey_hw1_problem3.py, e.g., I
22 will name my code as pratikac_hw1_problem3.py. Note, we will not accept .ipynb files (i.e.,
23 Jupyter notebooks), you should only upload .py files. If you are using Google Colab to do
24 your homework (and I suggest that you don’t...), you can export the notebook to a .py file.
25 • **In addition to submitting the code, you should append the entire Python code for the**
26 **particular problem to the solution in the PDF. If you are using Latex, you can do**
27 **something like the screenshot below. The instructors will execute the code to check it.**
28 **Your code should run without any errors and should create all output/plots required in**
29 **the problem.**

30

```
\includepackage{pythonhighlight}

\begin{python}

a = np.array(10)

...

\end{python}
```

- 31 **Credit** The points for the problems add up to 102. You only need to solve for 100 points to get full
32 credit, i.e., your final score will be $\min(\text{your total points}, 100)$.
-

33 **Problem 1 (20 points).** The torchvision library at <https://pytorch.org/vision/stable/models.html>
 34 implements a number of popular architectures that you can use quickly in your code. In this problem,
 35 you will take a deeper look at residual networks at
 36 <https://github.com/pytorch/vision/blob/master/torchvision/models/resnet.py>. Understand how this
 37 architecture is coded up.
 38 (a) (2 points) Do you notice peculiarities that you notice in the code, e.g., which one is better, using a
 39 batch-normalization layer before ReLU or after ReLU?
 40 (b) (2 points) What do the calls “model.train()” and “model.eval()” do and where do you use them in
 41 a typical training and validation code? Why did we not have them in HW 1 when we wrote our own
 42 library for training deep networks?
 43 (c) (3 points) Write code to iterate over all layers of the network and count the number of parameters
 44 in each layer of the network. Execute this code for the Resnet-18 network.
 45 (d) (3 points) Weight decay should not be applied to the biases of the different layers in the network,
 46 argue why this is the case.
 47 (e) (5 points) Write the code to iterate over all the network layers in Resnet-18 and separate out the
 48 parameters in three groups: (i) batch-norm affine transform parameters, (ii) biases of convolutional
 49 and fully-connected layers, and (iii) all the rest. There is no need to submit the code separately in this
 50 case, since these are a few lines of code just copy them out into your PDF solutions.
 51 (f) (5 points) Augmentations are increasingly becoming much more important for deep learning
 52 than what we initially believed. And therefore there are a number of sophisticated ways to per-
 53 form random augmentation. Go through the paper <https://arxiv.org/abs/1909.13719> (you don’t
 54 have to read the details in order to do this question) and its implementation in torchvision at
 55 <https://pytorch.org/vision/main/generated/torchvision.transforms.RandAugment.html>. The code of
 56 this implementation is at https://pytorch.org/vision/main/_modules/torchvision/transforms/autoaugment.html#RandAugment.
 57 Take any one image of your choice (you can use the same image of the astronaut as that of the
 58 examples) and show the result of augmenting this image using the following 10 augmentations: (a)
 59 ShearX, (b) ShearY, (c) TranslateX, (d) TranslateY, (e) Rotate, (d) Brightness, (e) Color, (f) Contrast,
 60 (g) Sharpness, (h) Posterize, (i) Solarize and (j) Equalize. You will use existing functions from
 61 torchvision to perform these augmentations (see the code of RandAugment linked here to understand
 62 how to call each of them). Each of these augmentations has some parameters that you will have
 63 to choose in order to run these corresponding augmentation functions. Any reasonable values are
 64 okay, feel free to experiment. The goal of this problem is to understand how these augmentations
 65 work. You don’t have to submit code in this case, pictures in the PDF correctly annotated with the
 66 parameters of the augmentation that was used to create them are fine.

67 **Problem 2 (12 points).** Non-convex optimization problems are harder than convex optimization
 68 problems. There are however a few special non-convex problems that are easy. We will look at one of
 69 them here, namely unconstrained matrix factorization. Given a matrix $X \in \mathbb{R}^{m \times n}$ we would like to
 70 decompose it into two matrices of rank at most r

$$X = AB$$

71 where $A \in \mathbb{R}^{m \times r}$ and $B \in \mathbb{R}^{r \times n}$. Think of arranging all your data as columns of X . Columns
 72 of the matrix A are like elements of a dictionary, they correspond to different patterns in the data
 73 and are called atoms. The matrix B chooses which patterns to collect together in order to create a
 74 particular datum, i.e., column of X . Solving for factors A, B is usually done with constraints, e.g., B
 75 is typically forced to be a sparse matrix which enables regenerating data X using as few atoms as
 76 possible. We will solve a simpler problem:

$$A^*, B^* = \underset{A \in \mathbb{R}^{m \times r}, B \in \mathbb{R}^{r \times n}}{\operatorname{argmin}} \|X - AB\|_F^2.$$

77 where $\|\cdot\|_F$ denotes the Frobenius norm.

78 (a) (2 points) Show that the above problem is not convex. You can give a counter-example.

79 (b) (6 points) The global optimum for this problem can be computed easily in spite of it being
 80 non-convex, find it. You may find it useful to write down the SVD of X .

81 (c) (4 points) Is the solution to the above optimization problem unique? Given one solution A^*, B^*
 82 name one way using which you can obtain another solution.

83 **Problem 3 (40 points). Use Google Colab to train the neural network, but after debugging
 84 everything on your laptop first.** Neural networks are high-dimensional classifiers. While we may be
 85 able to train and regularize SVMs to have a large margin between two classes, it is often difficult
 86 to do the same for neural networks. A consequence of this is that, roughly speaking, all samples in
 87 typical training datasets lie close to the decision boundaries after training. This makes it quite easy
 88 to make minor perturbations to the input image—perturbations that are imperceptible to the human
 89 eye—and send the sample across the decision boundary so as to cause the network to mis-predict.
 90 We will synthesize such adversarial perturbations in this problem. You can read more about it at
 91 <https://arxiv.org/abs/1412.6572>; however be wary of the heuristic generalizations in this paper that
 92 are incorrect.

93 We will find the best adversarial perturbation to a given image x and its target y , this amounts to
 94 solving the optimization problem

$$\max_{\|x' - x\|_\infty \leq \epsilon} \ell(x', y) \quad (1)$$

95 where x' is the *variable of optimization* and it is the adversarially perturbed image corresponding to x ,
 96 the quantity $\ell(x', y)$ is the loss computed on the image x' for the label y . We have chosen to make the
 97 parameters of the classifier w implicit for sake of clarity. This optimization problem searches for all
 98 images within an ϵ -ball of the original image x . We will use the ℓ_∞ -norm

$$\|x\|_\infty = \max_k |x_k|$$

99 in this problem. Let us perform the Taylor series expansion of the objective in Eq. (1)

$$\ell(x', y) = \ell(x, y) + \epsilon d^\top \nabla \ell(x, y) + \mathcal{O}(\epsilon^2);$$

100 here $d = (x' - x)/\epsilon$. We can now write an optimization problem for finding adversarial perturbations
 101 as

$$\max_{\|d\|_\infty \leq 1} d^\top \nabla \ell(x, y).$$

102 Notice that the constraint implies that any element of the vector d can be perturbed by at most 1; there
 103 is no limit on the number of elements perturbed. The value of d that maximizes this objective is

104 therefore the “signed gradient”

$$d_k = \frac{\nabla \ell(x, y)_k}{|\nabla \ell(x, y)_k|};$$

105 where $|\cdot|$ denotes the element-wise absolute value. If $\nabla \ell(x, y)_k < 0$, the corresponding $d_k < 0$ and
106 vice versa. The maximal objective is $\|\nabla \ell(x, y)\|_1$. The perturbation d is what we want to compute
107 next.

108 (a) (20 points) We will first train a convolutional neural network for this problem with all the bells
109 and whistles. You can use the code in allcnn.py provided at

110 <https://gist.github.com/pratikac/68d6d94e4739786798e90691fb1a581b>.

111 This is a small model with about 1.6M parameters. Train this model on the CIFAR-10 dataset for
112 100 epochs, you should try to get a validation error below 10%. You should use a GPU on Colab
113 for training; else your code will be very slow. Roughly speaking, running for 100–200 epochs will
114 take 2–4 hours, so be patient. You can use data augmentation such as mirror flips and brightness
115 and contrast changes to improve your validation accuracy. Plot the training and validation losses and
116 errors as a function of the number of epochs. Some hints for choosing hyper-parameters:

- 117 • Learning rate of 0.1 for the first 40 epochs, then 0.01 for the next 40 epochs and then 0.001
118 for the final 20 epochs.
- 119 • Weight decay of 10^{-3} .
- 120 • You will need to perform data augmentation, at least mirror flips, and cropping and padding.
- 121 • You will also need to use dropout and batch-normalization.
- 122 • Use SGD with Nesterov’s momentum of 0.9 to train the network.

123 Make sure you save the parameters of the network because we will need them for the next part.

124 (b) (10 points) We will next compute the backprop gradient of the loss with respect to the input. We
125 know that code of the form

```
126 ...
127
128 yh = net.forward(x)
129 loss = loss.forward(yh, y)
130
131 loss.backward()
```

133 computes the gradient of the loss with respect to the weights, i.e., it computes \bar{w} in our notation. You
134 can get the gradient \bar{x} very easily by adding the following line after `loss.backward()`.

```
135
136 dx = x.grad.data.clone()
```

138 Plot this gradient dx for a few input images which the network classifies correctly and also for a few
139 images which the network misclassifies. Comment on the similarities or the differences.

140 Note that each pixel of the RGB image x lies in $[0, 255]$, we will pick $\epsilon = 8$. Pick a particular
141 mini-batch $\{x_1, \dots, x_6\}$ with $b = 100$. For every image in this mini-batch perform the “5-step
142 signed gradient attack”, i.e., perturb that image 5 times using the signed gradient, at each step you
143 feed in the perturbed image from the previous step and perturb it a bit more. Your pseudo-code will
144 look as follows.

```

145 xs, ys = mini-batch of inputs and targets
146 for x,y in zip(xs, ys):
147     for k in range(5):
148         # forward propagate x through the network
149         # backprop the loss
150         dx = ...
151         x += eps*sign(dx)
152         # record loss on the perturbed image
153         ell = loss(x, y)
154

```

156 Plot the loss on the perturbed images as a function of the number of steps in the attack averaged
 157 across your mini-batch.

158 (c) (10 points) Compute the accuracy of the network on 1-step perturbed images, i.e., for every image
 159 in the validation set, perturb the image using a 1-step attack and check the prediction of the network.
 160 How does this accuracy on adversarially perturbed images compare with the accuracy on the clean
 161 validation set?

162 (d) (0 points) The human brain also has a lot of neurons and is likely a high-dimensional classifier.
 163 Are humans susceptible to adversarial perturbations? Can you give examples of images that fool the
 164 human visual system? Are these “small”, i.e., is $\|\epsilon/x\|_\infty$ small for these examples?

165 **Problem 4 (30 points, You can try to do this problem on your laptop when you are debugging but**
 166 **Colab with GPU will train much faster).** We will implement a model to predict the next character
 167 in a given sentence. We will use a few different sources of data to build this model (think carefully
 168 about how you will create a train and test set)

- 169 (1) The complete works of Shakespeare, which you can get at <https://www.gutenberg.org/ebooks/100.txt.utf-8> which is hosted at <https://www.gutenberg.org/ebooks/100>.
- 170 (2) Leo Tolstoy’s War and Peace as our training dataset. You can get the text file of
 171 the entire book from <https://www.gutenberg.org/ebooks/2600.txt.utf-8> which is hosted
 172 at <https://www.gutenberg.org/ebooks/2600>.
- 173 (3) Any English book of your choice which is sufficiently long, e.g., its text file is at least 1 MB.
 174 You can find text files for many books at Project Gutenberg <https://www.gutenberg.org>.

176 Many different ways of creating the train and test set are possible (they will give similar values for the
 177 eventual loss). We will not do cross-validation in this problem, so just create one train set and one test
 178 set (which is itself used as a validation set).

179 (a) (5 points) First observe that characters in the text can be letters, numbers, punctuation, and newlines.
 180 We will represent each character using its one-hot encoding; you should do

```

181 m = length(set(all_chars))
182

```

184 to get the correct number of unique characters. This is known as the size of the vocabulary in NLP.
 185 Do not worry about cleaning the dataset; neural networks are surprisingly good at handling unclean
 186 data. You will write a function that takes a part of the text input, say a sequence of 32 characters, and
 187 converts it into this embedding.

188 (b) (20 points) We now have written our problem in the standard form for an RNN where we are
189 given a long sequence of data x_1, x_2, \dots . Each element here is the embedding of a character or a
190 punctuation mark. You will train an RNN with one hidden layer (feel free to use a deeper network if
191 you'd like) which predicts the one-hot vector of the next *character* as output, given the past sequence.
192 This corresponds to the operations

$$\begin{aligned} h_{t+1} &= \tanh(w_h h_t + w_x x_{t+1} + b_h) \\ \mathbb{R}^{75} \ni \hat{y}_t &= \text{softmax}(w_y h_t + b_y) \\ \ell(\hat{y}_t, y_t) &= \ell(\hat{y}_t, x_{t+1}) = -\log(\hat{y}_t)_{x_{t+1}} \\ \ell &= \frac{1}{T} \sum_{t=1}^T \ell(\hat{y}_t, y_t). \end{aligned}$$

193 Notice that we are going to predict softmax output logits $\hat{y}_t \in \mathbb{R}^{75}$ over the entire vocabulary. The
194 loss is simply the cross-entropy loss of predicting the correct next character.

195 Code the RNN using PyTorch and train it. Plot the training loss and error of a mini-batch as a function
196 of the number of weight updates (you will see a noisy plot, that's okay). The validation loss/error of
197 an RNN is calculated the same way as the training loss/error, it is the cross-entropy loss on data that
198 was not a part of the training set. Report the validation loss and error over a future sequence of 32
199 characters (averaged over the size of the mini-batch), every 1000 weight updates.

200 (c) (5 points) You should also report a few examples of the RNN for generating new sentences. To do
201 so initialize the hidden state to zero and roll the RNN forward by setting x_{t+1} in the first equation
202 above to be the prediction of previous step \hat{y}_t . You will notice that although the RNN obviously does
203 not do a good job of generating correct words, it gets syntactic things like punctuation more or less
204 correct.

205 **Some tips:** You should use the inbuilt nn.RNN module to build the recurrent network. You should
206 also use torch.optim.Adam as the optimizer instead of SGD (we will study this soon) with a learning
207 rate of 10^{-3} . Set the length of the sequence in the mini-batch to be $T = 32$. You may have to clip the
208 gradients before calling optim.step() using the function torch.nn.utils.clip_grad_norm_. For more
209 tips read the code at https://github.com/pytorch/examples/tree/master/time_sequence_prediction and
210 https://github.com/pytorch/examples/blob/master/word_language_model/main.py carefully.