

UNIVERSITY OF PENNSYLVANIA

ESE 546: PRINCIPLES OF DEEP LEARNING

HOMEWORK 3

- You will submit solutions typeset in L^AT_EX on Gradescope (strongly encouraged). You can use hw_template.tex on Canvas in the “Homeworks” folder to do so. If your handwriting is unambiguously legible, you can submit PDF scans/tablet-created PDFs.
- Clearly indicate the name and Penn email ID of all your collaborators on your submitted solutions.
- Start a new problem on a fresh page and mark all the pages corresponding to each problem. Failure to do so may result in your work not graded completely.
- For each problem in the homework, you should mention the total amount of time you spent on it. This helps us keep track of which problems most students are finding difficult.
- You can be informal while typesetting the solutions, e.g., if you want to draw a picture feel free to draw it on paper clearly, click a picture and include it in your solution. Do not spend undue time on typesetting solutions.
- You will see an entry of the form “HW 1 PDF” where you will upload the PDF of your solutions. You will also see entries like “HW 1 Problem 1 Code” and “HW 1 Problem 3 Code” where you will upload your solution for the respective problems.
- **For each programming problem/sub-problem, you should create a fresh .py file.** This file should contain **all** the code to reproduce the results of the problem/sub-problem, e.g., it should save the plot that is required (correctly with all the axes, title and legend) as a PDF in the same directory. You will upload the .py file as your solution for “HW 1 Problem 3 Code” or “HW 1 Problem 3 Code”. Name your file as pennkey_hw1_problem3.py, e.g., I will name my code as pratikac_hw1_problem3.py. Note, we will not accept .ipynb files (i.e., Jupyter notebooks), you should only upload .py files. If you are using Google Colab to do your homework (and I suggest that you don’t...), you can export the notebook to a .py file.
- **In addition to submitting the code, you should append the entire Python code for the particular problem to the solution in the PDF. If you are using Latex, you can do something like the screenshot below. The instructors will execute the code to check it. Your code should run without any errors and should create all output/plots required in the problem.**

```
\includepackage{pythonhighlight}

\begin{python}

a = np.array(10)

...

\end{python}
```

Credit The points for the problems add up to 110. You need to solve for 100 points to get full credit, i.e., your final score will be $\min(\text{your total points}, 100)$.

1 **Problem 1 (35 points, do this problem on your laptop when you are debugging but Colab with**
 2 **GPU will train much faster).** The pre-training procedure for a large language model (LLM) involves
 3 minimizing the cross-entropy entropy loss for predicting the next token on a large corpus of data. In
 4 this problem, you will fit a small Transformer on the dataset that you used for problem 4 of HW 2 (a
 5 concatenation of the complete works of Shakesphere, Leo Tolstoy's War and Peace and any English
 6 book of your choice from Project Gutenberg). We wish to build a network that can complete a given
 7 sentence.
 8 (a) **(3 points)** First, you should tokenize the dataset using a more sophisticated tokenizer than
 9 the characters that you used last time. You will use what is called a byte pair encoding (BPE)
 10 https://en.wikipedia.org/wiki/Byte-pair_encoding.
 11 See the example shown at https://en.wikipedia.org/wiki/Byte-pair_encoding#Example. BPE works
 12 by scanning the text to identify frequently occurring byte pairs and collapsing them into a
 13 new symbol into the vocabulary. This is straight-forward to implement but it is quite slow in
 14 Python (it requires multiple passes over the dataset to compute the frequencies and compress the
 15 text). We will therefore use an existing library to build the vocabulary. Read the webpage at
 16 <https://huggingface.co/docs/tokenizers/quicktour> carefully to understand how to use Huggingface's
 17 tokenizers library (install it using "pip install tokenizers") to train and run a tokenizer.


18

```

1 From tokenizers import Tokenizer
2 From tokenizers.trainers import BpeTrainer
3 From tokenizers.models import BPE
4 From tokenizers.pre_tokenizers import BertPreTokenizer
5
6 tokenizer = Tokenizer(BPE(unk_token="[UNK"]))
7 tokenizer.pre_tokenizer=BertPreTokenizer()
8
9 trainer = BpeTrainer(special_tokens=[ "[UNK]", "[CLS]", "[SEP]", "[PAD]", "[MASK]" ])
10 tokenizer.train(['enwiki18.txt'],trainer)

```

19 You can set up your code to do something like the picture above. Two things to note here. First,
 20 we have showed an example that uses a few special tokens while tokenizing the text, e.g., [UNK]
 21 stands for “unknown token” for when the data might contain new characters that were not a part of the
 22 vocabulary, the [CLS] token is sometimes used for downstream tasks using pre-trained BERT/GPT
 23 representations, [PAD] stands for padding which allows the user to create a mini-batch of short and
 24 long sentences, etc. The vocabulary of chatbots contains many such special tokens for start of turn,
 25 end of turn, beginning of “thinking”, end of “thinking” etc. Second, we are using a pre-tokenizer
 26 from Bert, this procedure simply sanitizes the text, e.g., splitting apostrophes with a space on both
 27 sides (so that the tokenizer will have the apostrophe symbol as an element of the vocabulary), etc.
 28 Such procedures are broadly called stemming and lemmatization in the NLP literature and they are a
 29 very important step before any language modeling can be done (Spacy <https://spacy.io/api/tokenizer>
 30 is an excellent library for this purpose).
 31 You can give options to BpeTrainer to use a particular vocabulary size, use something in the region of
 32 1024–2048.
 33 (b) **(17 points)** Next you will write code for a (decoder-only) Transformer-based network. You
 34 may not use pre-coded Transformer blocks from any library (this includes PyTorch layers such as
 35 nn.Transformer, nn.TransformerEncoder etc.). You should write the following code by yourself using
 36 basic PyTorch tensors. Here is a rough guideline for the various parameters is as follows.



37

- A context length of $T = 128$.
- Each token is embedded as a vector into $d = 256$ dimensions with sine-cosine-based position encoding.
- The feature dimension of keys, queries and values is $p = 32$, the number of heads is 2.
- Each Transformer block has one self-attention-based layer and one MLP layer, with layer normalization and residual connections as shown in the picture above. We will code up all these blocks ourselves. You can set the bias to False for all operations (linear operations, key/value computation, as well as normalization).
- For multi-head self-attention if input features are $h \in \mathbb{R}^{T \times p}$, the computations of the multi-head attention block for each head denoted by the index $m = 1, 2$ are:

$$\mathbb{R}^{T \times p} \ni k^m = \sigma(h w_k^m) \text{ with } w_k^m \in \mathbb{R}^{p \times p}$$

$$\mathbb{R}^{T \times p} \ni q^m = \sigma(h w_q^m)$$

$$\mathbb{R}^{T \times p} \ni v^m = \sigma(h w_v^m)$$

$$\forall t : \mathbb{R}^p \ni h'^m_t = \text{dropout} \left(\sum_{s=1}^{t-1} \left(\frac{\exp(k_s^m \top q_t^m / \sqrt{p})}{\sum_{u=1}^{t-1} \exp(k_u^m \top q_t^m / \sqrt{p})} \right) v_t^m \right).$$

48 We then concatenate the output of each head to get features $h' \in \mathbb{R}^{T \times 2p}$ before an output
49 projection layer

$$\mathbb{R}^{T \times p} \ni h_{\text{attn}} = h' w_o \text{ with } w_o \in \mathbb{R}^{2p \times p}.$$

- 50 • Residual connection and normalization are easy to write

$$\mathbb{R}^{T \times p} \ni h_{\text{attn}} = \text{layer norm}(h + h_{\text{attn}})$$

51 You can use PyTorch's nn.LayerNorm to implement this operation.

- 52 • The next thing we should is to implement a position-wise MLP to mix the features computed
53 by self-attention. This looks like

$$\mathbb{R}^{T \times p} \ni h = \text{layer norm}(h_{\text{attn}}(t) + \text{MLP}(h_{\text{attn}}(t)))$$

54 with the MLP being a two-layer network with ReLU nonlinearities, p input neurons, $2p$ hidden
55 neurons and p output neurons. Notice that the MLP acts on the input data position-wise.

- 56 • All this together forms one “block” of a Transformer. Your network should have 2–4 such
57 blocks.

58 (c) **(10 points)** Train this network on the tokenized data for predicting the next token using the past
59 tokens in the context (you will see that the equations above have causal attention). In this problem, we
60 are only interested in completing a given sentence, so we will create a mini-batch where each datum
61 is a sentence. Remember that you will need to use padding while creating mini-batches that contain
62 sentences of unequal lengths.

63 You should plot the (i) training loss of successive mini-batches as a function of the weight updates
64 (feel free to do a moving average), (ii) periodically, say after every 1000 weight updates, you can also
65 compute the validation loss on the validation dataset.

66 (d) **(5 points)** Write code for auto-regressive sampling using this network. Given a prompt, say “The
67 Eiffel Tower is in Paris”, your code will tokenize the prompt (let’s say this prompt contains 6 tokens
68 x_1, x_2, \dots, x_6), and run the sampler step-by-step as


$$\begin{aligned}\hat{x}_7 &\sim p(\cdot | x_1, \dots, x_6) \\ \hat{x}_8 &\sim p(\cdot | x_1, \dots, x_6, \hat{x}_7) \\ \hat{x}_9 &\sim p(\cdot | x_1, \dots, x_6, \hat{x}_7, \hat{x}_8),\end{aligned}$$

69 each sample involves a call to the network and a random draw from the probability distribution of the
70 logits. Remember that you need to use position-encoding when you feed-in the context. You should
71 report a few examples of the network completing your prompts (use prompts of length 5–6 words,
72 completions of length 7–8 words).

73 **Problem 2 (20 points. Do this problem on your laptop; you do not really need a GPU for this
74 problem.).** In this problem, we will use nonlinear regression to understand one possible reason why
75 deep networks can avoid overfitting even when the number of samples in the dataset is smaller than
76 the number of weights. We will effectively show that generalization occurs when the test data lies
77 within the “convex hull” of the training data. If the test samples lie far away from the train data, then
78 one should not expect any generalization with networks with lots of parameters such as deep networks.
79 A second purpose of working on this problem is to see that we can do interesting experiments and
80 train 1000s of networks on a laptop (even without a GPU).

81 (a) **(0 points)** You will create the dataset yourself. Let the true data come from

$$\mathbb{R} \ni y = f^*(x) \equiv \sin(10\pi x^4); \text{ for } x \in [0, 1].$$



83 You will write a function to sample n inputs $x_i \in [0, 1]$ and their corresponding outputs $y_i =$
84 $\sin(10\pi x_i^4)$ to create a dataset $D_n = \{(x_i, y_i)\}_{i=1}^n$.

85 (b) (10 points) Build a multi-layer perceptron (MLP) in PyTorch to regress the dataset. You can use
86 an MLP with 2-3 layers, ReLU nonlinearities with batch-normalization after each layer and train it
87 with SGD. You can use any PyTorch function for this that you would like; but it is actually easier to
88 code the mini-batch sampling etc. yourself simply using a for loop for the iterations instead of using a
89 dataloader. You should choose the hyper-parameters (weight-decay, learning rate, number of epochs)
90 to ensure that you are getting zero or very close to zero training error. Denote the learned function as

$$\hat{y} = f_w(x; n)$$

91 where w are the weights of the MLP, x is the test input and we have denoted by ; n the fact that this
92 model was trained with n data points. Make sure your code for training the MLP is efficient, we will
93 be fitting about 100 different models. Solving this problem will also convince you that your laptops
94 are plenty powerful to ask serious questions about why deep networks work; you do not need GPUs
95 for everything in deep learning.

96 Write a function to evaluate the quantity

$$\delta f_{\text{in}}(n) = \max_{x \in [0, 1]} |f_w(x; n) - f^*(x)|$$
$$\delta f_{\text{out}}(n) = \max_{x \in [0, 1.5]} |f_w(x; n) - f^*(x)|.$$

97 You will evaluate the maximum in the above expression by sampling many (say 1000) test inputs
98 $x \in [0, 1]$ or $x \in [0, 1.5]$ and taking the maximum of their predictions. The quantity $\delta f_{\text{in}}(n)$ is the
99 largest discrepancy within the support of the dataset, i.e., within $[0, 1]$, between the true function
100 $f^*(x)$ and our fitted function $f_w(x; n)$ using n training data points. Similarly the quantity $\delta f_{\text{out}}(n)$ is
101 the largest discrepancy outside the support, for example in $x \in [0, 1.5]$.

102 (c) (10 points) For each of 20 different values of n , e.g., $ns = np.logspace(1, 3, 20).astype(int)$, create
103 5 training datasets (i.e., you will create 100 different datasets in total) like it is described in part
104 (a), fit an MLP to this dataset and evaluate the mean and standard deviation (across the 5 datasets
105 corresponding to each n) of both $\delta f_{\text{in}}(n)$ and $\delta f_{\text{out}}(n)$. Draw a plot of $\log \delta f_{\text{in}}(n)$ (mean and standard
106 deviation) versus n . On the same plot also draw $\log \delta f_{\text{out}}(n)$ (mean and standard deviation) versus n .
107 Interpret your plot.

108 **Problem 3 (20 points).** In this problem you will show that co-coercivity of the gradient and its
109 Lipschitz continuity are equivalent. Assume that the function $f(x)$ is convex.

110 (a) (5 points) For a differentiable function $f : \mathbb{R}^d \rightarrow \mathbb{R}$, show that co-coercivity of ∇f implies
111 Lipschitz continuity of ∇f . In other words, show that for all x, y

$$\langle \nabla f(x) - \nabla f(y), x - y \rangle \geq \frac{1}{L} \|\nabla f(x) - \nabla f(y)\|^2 \implies \|\nabla f(x) - \nabla f(y)\| \leq L \|x - y\|.$$

112 Hint: use the Cauchy-Schwartz inequality.

113 (b) (8 points) Show the converse, i.e., show that Lipschitz continuity of ∇f implies co-coercivity.

114 (c) (7 points) For a twice-differentiable function $f(x)$ with L -Lipschitz gradients and strong-convexity
115 parameter m , show that

$$m \leq \|\nabla^2 f(x)\|_2 \leq L$$

116 for all x . Hint: use the mean-value theorem on a line that joins two points x, y .

117 **Problem 4 (35 points). (Do this on your laptop)** In this problem, we will implement logistic
118 regression for classifying two classes (zero and one) from MNIST. You may not use any routines from
119 PyTorch other than the ones that help download the data. Use Numpy to code up the optimization
120 algorithm.

121 (a) **(0 points)** First, prepare the dataset. Select all the samples belonging to the first two classes from
122 MNIST's training dataset, this will be your training dataset. Similarly, create a validation dataset for
123 the two classes by picking samples corresponding to the first two classes from the validation set of the
124 MNIST dataset. You can subsample input images from 28×28 to 14×14 if you need.

125 (b) **(10 points)** Logistic regression solves for

$$\underset{w \in \mathbb{R}^d, w_0 \in \mathbb{R}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y_i (w^\top x_i + w_0)} \right) + \frac{\lambda}{2} \left(\|w\|_2^2 + w_0^2 \right) \quad (1)$$

126 where $x \in \mathbb{R}^{196}$. Set $y_i = 1$ for MNIST images labeled zero and $y_i = -1$ for MNIST images labeled
127 one. Initialize the weights randomly for both the following parts but make sure that they are the same
128 for both gradient descent and gradient descent with Nesterov's acceleration in part (c). You can try a
129 few different values of λ and pick the one that gives the best validation error for the following parts.

130 Optimize the objective in Eq. (1) using gradient descent (note, not stochastic gradient descent) and
131 plot the training loss as a function of the number of parameter updates on a semi-log scale (log scale
132 on the Y-axis). This plot should be a straight line. As we saw in the class, the slope of this line should
133 be about $-\kappa^{-1}$ for gradient descent. Compute the slope of the line in your plot and mention it clearly.

134 (c) **(5 points)** Write down the Hessian of the loss function in Eq. (1). Without assuming any special
135 conditions about the dataset $\{(x_i, y_i)\}_{i=1, \dots, n}$, is this problem strongly convex? What is the best
136 strong convexity parameter for the loss function in Eq. (1)?

137 (d) **(10 points)** Optimize the objective in Eq. (1) again, this time using gradient descent with Nesterov's
138 acceleration. If we knew the condition number κ , what momentum coefficient would we use for this
139 problem? It is difficult, although possible, to evaluate κ , so let's use

$$\kappa = \frac{L}{m}$$

140 where we treat L as a hyper-parameter, i.e., we choose a value for L and set m to be the solution of
141 part (b). Again, choose the value of L by trying out a few values and looking at the slope of the
142 training loss for each setting. (Hint: the momentum parameter is typically between 0.75-0.95).

143 The slope of the semi-log plot of training loss versus the number of parameter updates in this case
144 will be about $-\kappa^{-1/2}$ for Nesterov's updates. Note that we do not know the correct κ in this problem,
145 so your slope may not match the value you chose for κ above. You should however see that the slope
146 of the plot for Nesterov's acceleration is better than the slope of the plot you obtained in part (b).

147 (e) **(10 points)** We will now optimize the same problem with stochastic gradient descent (SGD). Use
148 a batch-size $b = 128$ (feel free to try different batch-size such as 64 and 8) and optimize Eq. (1) using
149 SGD with and without Nesterov's acceleration. Plot the training loss against the number of parameter
150 updates on a semi-log scale (log scale on the Y-axis) for (i) gradient descent with Nesterov's updates
151 (can be the same plot from your previous solutions), (ii) SGD without Nesterov's acceleration and,

¹⁵² (iii) SGD with Nesterov's acceleration. Is the convergence for (iii) faster than that of (ii)? Comment
¹⁵³ on the differences for the convergence curves of (i) and (ii).